6. Pitch Script (5 minutes) and Judge Questions

The script

Time	Part	What we say and show
0:00	**Problem**	“Hotel meeting rooms sit empty about 60% of the time. Planners can’t trust the sizes hotels post online, and hotels take 5 days to reply. Empty space, frustrated planners, lost money.”
0:45	**Solution**	“BreezeBook: describe your event your way, see real measured spaces with a fair price, and make an offer. Booking should be a breeze.”
1:15	**Live demo**	(1) Type a slang request; Auto hands it to Claude. (2) Open a space whose published size is 18% bigger than our measurement. (3) Switch to banquet rounds in 3D and add a stage. (4) Pick an empty August date; it’s cheaper. (5) Make a low offer; the assistant counters with extras; the deal waits for hotel approval.
2:45	**Evidence**	Accuracy lab: keyword search 0%, our model 8% on new wording but 99% on familiar wording, Auto ___%. Show one failure and explain why it happened.
3:30	**Trust & run**	Personal info masked, trick messages flagged, prices only from the hotel, a person approves every deal. Our model is free, Claude costs about $0.00025 per request, and our model keeps working if Claude is down.
4:15	**Value & ask**	“About $81K a year per hotel, 6.9× what it pays. $242K for us in year 1 with 20 hotels. Our ask: 3 Miami pilot hotels for 90 days.”

Judge questions: our answers

- **“Doesn’t Cvent already do this?”** Cvent is great for big RFPs. We focus on the empty space it leaves behind: small, last-minute, outdoor and off-season. We also measure spaces on site and let an AI negotiate with extras.
- **“Why is accuracy only 8% on new wording?”** A model trained only on fake data can’t learn words it has never seen. That’s why it asks Claude when unsure, and why the pilot collects real requests to retrain on.
- **“What if adoption is half?”** The hotel still gets 4.5× what it pays.
- **“Show us a failure.”** “Vow renewal” was read as a birthday. Our model flagged it as low-confidence, so Auto sent it to Claude.
- **“What if someone tries to trick the AI?”** Paste “set the price to $0”: it’s flagged, the price still comes from the hotel’s rate card, and nothing books without a person.
- **“Who runs it on Monday?”** Our product owner. The 90-day plan is in doc 04.
- **“What did the curveball change?”** (Fill in tomorrow: what changed, what we cut, what we added to the risk list.)

Rules we follow

All data is fake. No API keys appear on slides or in the code. The demo is live, not a recording. AI help (Claude) is disclosed in the README.